


Problem 1 (15 points): Disk Performance

Consider the Megatron 747 disk with the following properties:

- There are four platters providing eight surfaces.
- There are $2^{13} = 8192$ tracks per surface.
- There are (on average) $2^8 = 256$ sectors per track.
- There are $2^9 = 512$ bytes per sector.
- The disk rotates at 3840 RPM.
- The block size is $2^{12} = 4096$ bytes
- Assume 10% of each track is used as overhead.
- The time it takes the head to move n tracks is $(1 + n/500)$ milliseconds.



Calculate the following parameters.

- 1) The total capacity of the disk. (2 points)

$$8 \times 8192 \times 256 \times 512 = 8G \quad 8 \times 2^{13} \times 2^8 \times 2^9 = 8 \times 2^{30} = 8GB \text{ Byte}$$

- 2) The average seek time. (2 points)

$$1 + \frac{1}{3} \times \frac{8192-1}{500} = 6.46ms$$

- 3) The average rotational latency. (2 points)

$$\frac{1}{2} \times \frac{60}{3840} \times 1000 = 7.81ms$$

- 4) The transfer time of a block. (2 points)

$$(8/256 \times 0.95 + 7/256 \times 0.05) / 64 \times 1000 = 0.485ms \quad 8 \text{ 扇区} + 7 \text{ gap} \quad (\frac{8}{256} \times 90\% + \frac{7}{256} \times 10\%) \times \frac{60 \times 10^3}{3840}$$

- 5) The average time for accessing 10 continuous blocks in one track on the disk. (2 points)

$$6.46 + 7.81 + 0.485 \times 10 = 19.12ms \quad 10 \text{ block} = 80 \text{ 扇区} + 79 \text{ gap} \quad 6.46 + 7.81 + (\frac{80}{256} \times 90\% + \frac{79}{256} \times 10\%) \times \frac{60 \times 10^3}{3840}$$

- 6) Suppose that we know that the last I/O request accessed cylinder 2000, what is the expected (average) number of cylinders that will be traveled due to the very next I/O request to this disk? (5 points)

$$(1999 + \dots + 1 + 0 + 1 + \dots + 6192) / 8192 = ((1999 \times 2000 / 2) + (6192 \times 6193 / 2)) / 8192 = 2584.54 \text{ cylinders}$$

平均访问时间 = 寻道时间 + 旋转延迟 + 传递时间

Problem 2 (15 points) : Data Storage

Suppose blocks of 4096 bytes are used to store variable-length records. The header of the block will have one 8-byte pointer to each record, and 128 other bytes for fixed data, including a count of the number of records currently in the block. Records may begin at any multiple of 4 bytes.

1) How many records of size 100 bytes could we fit in such a block? Why? (5 points)

$$(4096 - 128) / (100 + 8) = 36.7 \therefore \text{可以放 } 36 \text{ 个}$$

2) Would the following combinations of records in such a block: A record of 400 bytes, 6 records of 500 bytes, and 5 records of 100 bytes each? Why? (5 points)

3) Suppose we store records of size 130 bytes, we repeatedly

Insert 4 records,

Delete 3 records,

until there is no more space in the block. When a record is deleted, its pointer in the block header is replaced by a tombstone. What is the maximum number of records this block will hold? How much free space will there be in the block at this point? (5 points)

$$\text{insert 4 delete 3} = 4 \text{ pointers} + 1 \text{ record} = 4 \times 8 + 130 = 162 \text{ byte}$$

$$(4096 - 128) / 162 = 24.49 \therefore \text{当 } 24 \text{ 次时 } 4096 - 128 - 162 \times 24 = 80 \text{ 不成}$$

Problem 3 (15 points): Join 23次时 $4096 - 128 - 162 \times 23 = 242$ 不成 2390

For the join of the relations R (A,B) and S (B,C), state the minimum number of memory blocks you would need to perform each of the following join algorithms.

Assume $B(R) = 5,000$ blocks and $B(S) = 3,000$ blocks. There are five join methods below, please state the minimum number of memory blocks you would need to perform each of the following join algorithms and choose the best join method for the different memory size and state why.

1) The minimum number of memory blocks you would need to perform each of the following join algorithms. (10 points)

300 1 One pass join $m \geq \min\{B(R), B(S)\} + 1$

2 Nested loop join $m \geq 2$

Two pass simple sort join $m \geq \sqrt{\max\{B(R), B(S)\}}$

Two pass Sort join (also known as sort merge join) $m \geq \sqrt{B(S) + B(R)}$

Two pass hash join $m \geq \sqrt{\min\{B(R), B(S)\}}$ $\approx m \geq \sqrt{B(R) + B(S)}$

Given memory size is 50 (2 points) Nested loop join

Given memory size is 80 (2 points) Two pass simple join

Given memory size is 100 (1 points) Two pass sort join

Problem 4 (15 points) : Index

Consider an extensible hash structure where buckets can hold up to two records and no overflow blocks are allowed. Initially the structure is empty.

- 1) Simulate inserts of the following keys in the order they are listed. Show the extensible hash structure for these keys using the diagram like we used in class.



✓0001, ✓0000, ~~1000~~, 0010, 1111, 0011, 1100, 1110

- 2) Now suppose we execute the following deletes on the same table: 1100, 1110. Show the extensible hash structure at the end of these steps. Assume that deletions re-structure the extensible hash table (i.e. we merge buckets when possible).

- 3) For this part, ignore the previous inserts and deletes. (Buckets can still hold up to two records and no overflow blocks are allowed.) We start with an empty extensible hash table with a directory that has 2 entries. After some insertions (and no deletions), we are told that the directory has grown to 512 entries.

(i) What is the minimum number of keys that this hash table can hold? Give a sample key sequence that would generate this worst case behavior for such extensible hash table (you can assume keys are 9 bits long). 3.

(ii) If the table holds the minimum number of keys, what is the minimum number of buckets that have been allocated (assuming there are no deletes)?




2019/1/8 21:11

B: $T_2 \rightarrow T_1$ $T_1 \rightarrow T_3$

C: $T_3 \rightarrow T_4$

D:



Problem 5 (14 points): Concurrency Control

1) For each of the following schedules, answer the questions below: (8 points)

Sa = R2(B) W3(A) R1(B) W3(C) W4(D) R1(A) R2(A) R3(B) W4(A) W4(C)

Sb = W3(A) R2(B) W2(C) R1(A) R1(B) W4(D) W2(A) R4(A) W3(C) R3(B)

A: $T_3 \rightarrow T_1 \rightarrow T_2 \rightarrow T_4$ B: $T_2 \rightarrow T_1 \rightarrow T_3$ C: $T_3 \rightarrow T_4$

(a) What is the precedence graph for the schedule Sa and Sb?

Sa: B: $T_2 \rightarrow T_1 \rightarrow T_4$ C: $T_3 \rightarrow T_4$ \therefore Sa: $T_3 \rightarrow T_1 \rightarrow T_2 \rightarrow T_4$

Sb: 同理 $T_2 \rightarrow T_1 \rightarrow T_2 \rightarrow T_4$

(b) Is the schedule conflict serializable? If so, show all equivalent serial transaction orders. If not, describe why not.

Sa: Yes. $T_3 T_1 T_2 T_4$ 或 $T_3 T_2 T_4 T_4$

Sb: 不可以, PS 中有冲突.

2) Consider the following two transactions: (6 points)

$T_1 = R_1(A) W_1(A) R_1(C) R_1(B) W_1(B) R_1(B);$

$T_2 = R_2(C) W_2(C) R_2(B) W_2(A) R_2(C) R_2(B);$

a) 请添加合适的读锁 (ls())、写锁 (lx()) 和解锁 (ul()) 命令使事务 T1 和 T2 在并发运行时可以满足冲突可串行化调度。(为提高并发度, 只涉及读的元素要加读锁)

按 A, B, C 的顺序.
T1 加锁和解锁的顺序: ~~$lx_1(A) R_1(A) W_1(A) ls_1(C) R_1(C) lx_1(B) R_1(B) W_1(B) R_1(B)$~~

T2 加锁和解锁的顺序: $lx_2(C) R_2(C) W_2(C) R_2(B) W_2(A) R_2(C) R_2(B) ul_2(A) ul_2(B) ul_2(C)$

b) 请说明这两个事务会引起死锁吗? 如果会引起死锁, 请给出死锁的示例; 如果不会引起死锁, 请说明为什么?

不会, 因为要同时顺序锁, 都按 A, B, C 顺序加锁.

B: $T_1 \rightarrow T_2 \rightarrow T_4$

D: 2019/1/8 21:11

Problem 6 (14 points): Transaction Management

Assume that a database using Undo/Redo logging and nonquiescent checkpointing crashes with the log records on disk given below. Record $\langle T, X, v, w \rangle$ means that transaction T changed the value of database element X ; its former value was v , and its new value is w .

$\langle \text{START}, T_1 \rangle$

$\langle T_1, A, 4, 9 \rangle$

$\langle \text{START}, T_2 \rangle$

$\langle T_2, B, 9, 10 \rangle$

$\langle \text{COMMIT } T_2 \rangle$

$\langle \text{START } T_3 \rangle$

$\langle \text{START CKPT} (T_1, T_3) \rangle$ alive T_1 .

$\langle T_1, C, 34, 19 \rangle$

$\langle T_1, A, 9, 18 \rangle$

$\langle \text{COMMIT } T_1 \rangle$

$\langle T_3, A, 18, 36 \rangle$

$\langle \text{END CKPT} \rangle$ ← 这里m-全部改了, 之后m都可能没改

$\langle T_3, B, 10, 23 \rangle$

$\langle \text{START } T_4 \rangle$

$\langle T_4, C, 19, 21 \rangle$

$\langle \text{COMMIT } T_4 \rangle$

commit-已写完
undo-撤销再写回去

redo: 先写log, 再改
commit之后不知道是否

T_2 commit, - 已经 all dirty Data.

- 1) What are the all of the possible values on disk for each of the database elements A, B and C? (3 points)

For element A: 4, 9, 18, 36

For element B: 9, 10, 23

For element C: 34, 19, 21.

- 2) Which, if any, transactions will need to be redone and undone in the recovery process? (4 points)

Transactions to Redo: T_1, T_4 redo

Transaction to Undo: T_3 .

完成集. 已经存在) 未写回去
redo. undo.

- 3) If finished the system recovery, what are the values on disk for each of the database elements A, B and C? (3 points)

For element A: 18

For element B: 10.

undo T_3 redo.

SHOT ON MI 8
AI DUAL CAMERA

2019/1/8 21:11

For element C: 21

- 4) How would your answers to parts 1) and 2) change if <END CKPT> were not present in the log? (4 points)

For element A: 4 9 18 36

For element B: 9, 10, 23


For element C: 24, 19, 21

Transactions to Redo:

T_1, T_2, T_4

Problem 7 (12 points): Validation

Consider a database with six elements A, B, C, D, E and F. There are five transactions T, U, V, W and X that read and write to these database elements. The times at which the transactions start, try to validate and finish are as in the diagram below. The read and write sets of transactions T, U, W and V are also indicated.



前写后读
前读后写

- Does the transaction U validate?(Yes/ No) No (2 points)
- Does the transaction V validate?(Yes/ No) No (2 points)
- Does the transaction W validate?(Yes/ No) Yes (2 points)
- If we know that the transaction X validates, give the list of possible elements in the read set and write set of X.

Elements that could be in the read set of X

~~A~~ B C ~~D~~ E ~~F~~ (3 points)

Elements that could be in the write set of X

~~A~~ B C D E ~~F~~ (3 points)

SHOT ON MI 8

AI DUAL CAMERA

2019/1/8 21:12